

GateBreaker: Gate-Guided Attacks on Mixture-of-Expert LLMs

Lichao Wu

*System Security Lab,
Technical University of Darmstadt*

Sasha Behrouzi

*System Security Lab,
Technical University of Darmstadt*

Mohamadreza Rostami

*System Security Lab,
Technical University of Darmstadt*

Stjepan Picek

*Faculty of Electrical Engineering and Computing,
University of Zagreb & Radboud University*

Ahmad-Reza Sadeghi

*System Security Lab,
Technical University of Darmstadt*

Abstract

Mixture-of-Experts (MoE) architectures have advanced the scaling of Large Language Models (LLMs) by activating only a sparse subset of parameters per input, enabling state-of-the-art performance with reduced computational cost. As these models are increasingly deployed in critical domains, understanding and strengthening their alignment mechanisms is essential to prevent harmful outputs. However, existing LLM safety research has focused almost exclusively on dense architectures, leaving the unique safety properties of MoEs largely unexamined. The modular, sparsely-activated design of MoEs suggests that safety mechanisms may operate differently than in dense models, raising questions about their robustness.

In this paper, we present GateBreaker, the first training-free, lightweight, and architecture-agnostic attack framework that compromises the safety alignment of modern MoE LLMs at inference time. GateBreaker operates in three stages: (i) *gate-level profiling*, which identifies safety experts disproportionately routed on harmful inputs, (ii) *expert-level localization*, which localizes the safety structure within safety experts, and (iii) *targeted safety removal*, which disables the identified safety structure to compromise the safety alignment. Our study shows that MoE safety concentrates within a small subset of neurons coordinated by sparse routing. Selective disabling of these neurons, approximately 3% of neurons in the targeted expert layers, significantly increases the averaged attack success rate (ASR) from 7.4% to 64.9% against the eight latest aligned MoE LLMs with limited utility degradation. These safety neurons transfer across models within the same family, raising ASR from 17.9% to 67.7% with one-shot transfer attack. Furthermore, GateBreaker generalizes to five MoE vision language models (VLMs) with 60.9% ASR on unsafe image inputs. To our knowledge, no prior work achieves this level of efficacy against MoE LLMs.

1 Introduction

Large Language Models (LLMs) have achieved remarkable advances in natural language processing, powering applica-

tions ranging from search engines and coding assistants to scientific discovery and healthcare decision-making [1–4]. Their rapid adoption is largely driven by their ability to generalize across diverse tasks, often outperforming specialized systems. Scaling laws have played a crucial role in advancing LLMs, showing that model performance continues to improve with increasing data and parameters [5]. Yet, dense scaling incurs prohibitive computational and memory costs. Mixture-of-Experts (MoE) architectures have emerged as a compelling alternative, achieving state-of-the-art performance while maintaining computational efficiency [6, 7]. By activating only a sparse subset of experts per input, MoE LLMs enable trillion-parameter models to operate at a fraction of the cost of dense architectures. Recent breakthroughs, such as GPT-4 [8], DeepSeek-MoE [9], and Alibaba’s Qwen-MoE [10, 11], highlight the central role of MoEs in the next generation of LLMs. Architecture variants, such as sparse [11–14], mixture [9, 10, 15], and grouped mixture [16], are actively proposed and adopted in the latest LLM models and applications.

Attack on LLMs. While LLMs achieve significant success in various domains, it also amplifies concerns about safety and reliability: when misused or misaligned, LLMs can generate harmful, deceptive, or biased outputs [17, 18]. To mitigate LLM unsafe behaviors, techniques such as reinforcement learning from human feedback (RLHF) [19], direct preference optimization [20], and red teaming [21] have been adopted to align LLM with human values, denoted as *safety alignment*. Unfortunately, attacks on LLMs, especially on dense architectures, have revealed significant weaknesses in current alignment methods. Existing methods, such as adversarial prompting [22, 23], model editing [24, 25], and lightweight fine-tuning [26], could effectively bypass the safety alignment in the dense LLMs.

While safety challenges are already pressing in dense LLMs, they are likely amplified in MoE architectures. Unlike dense models, which activate all parameters for every input, MoE LLMs engage only a small subset of experts per token through a learned gating mechanism. Specifically, each input token passes through a small feed-forward module, known as a gate

(or router), that scores all available experts and selects the top- k experts to activate based on the token’s content. This conditional computation introduces new safety dynamics: different inputs may route to disjoint sets of experts, meaning that safety alignment is no longer uniformly distributed across the model but instead concentrated within a subset of experts. As a result, harmful prompts may bypass safety mechanisms entirely if routed to inadequately aligned experts. Moreover, the sparse and input-dependent nature of expert activation makes it difficult to ensure consistent and robust refusal behavior, as the pathways for safety alignment may be fragile or incomplete. These architectural traits create novel and realistic failure modes. For example, a seemingly aligned MoE LLM-based assistant could be manipulated into generating misinformation in healthcare or finance simply by steering prompts toward misaligned experts. In content moderation systems, adversaries could craft inputs that avoid triggering safety experts, allowing toxic or policy-violating content to pass undetected. Similarly, jailbreak prompts could exploit fragmented alignment pathways to elicit malicious behaviors that would otherwise be blocked in dense models. Despite these concerns, existing research remains limited: prior works operate only at coarse expert granularity and focus primarily on sparse experts [27, 28], leaving the fundamental safety behavior and corresponding weaknesses across diverse, state-of-the-art MoE designs largely unexplored. Given the growing centrality of MoEs in modern LLMs, there is an urgent need for fine-grained investigations of their safety properties.

Our Contribution. In this paper, we propose GateBreaker, the first training-free and lightweight attack framework that generalizes to state-of-the-art MoE LLMs with different architectures. GateBreaker operates in three stages: *gate-level profiling* to identify safety-relevant experts disproportionately triggered by harmful inputs. Guided by the gate profiling, we perform in-depth *expert-level localization* to localize safety-relevant neurons in selected local expert layers. Finally, we compromise safety alignment with *targeted safety removal*. We demonstrate that MoE safety behavior emerges from a small fraction of neurons spread across a subset of experts, coordinated through sparse routing patterns; simply masking them is sufficient to significantly increase the attack success rate (ASR), with negligible utility degradation. Concretely, our contributions are:

- We present a comprehensive analysis of safety alignment in MoE LLMs, revealing how safety behaviors are structurally distributed across experts and layers.
- We propose GateBreaker, the first attack framework designed to compromise safety alignment in MoE LLMs. GateBreaker is architecture-agnostic and generalizes across diverse open-weight MoE variants from leading developers, including OpenAI, DeepSeek, and Alibaba.
- We develop a lightweight, inference-time pipeline that performs gate-level profiling, expert-level localization,

and targeted neuron-level pruning in sequence. This modular design enables precise safety interventions with minimal overhead. The entire process executes during inference time, making it broadly accessible and suitable even for low-resource adversaries.

- GateBreaker achieves a state-of-the-art ASR. Across eight state-of-the-art reasoning and non-reasoning MoE LLMs, our method yields an average ASR of 64.9%, with fewer than 2.9% of neurons modified per targeted layer. Utility evaluations show negligible degradation on benign tasks. Additionally, GateBreaker generalizes to MoE vision language models, increasing ASR from 20.8% to 60.9% across five models using unsafe images.
- GateBreaker demonstrates strong cross-model transferability. Safety neurons identified in one model can be reused to attack sibling variants, such as domain- or instruction-tuned derivatives. In one-shot transfer settings, we observe ASR gains of up to 91.2%, exposing shared structural vulnerabilities in safety alignment across the MoE model family.

The remainder of this paper is structured as follows. Section 2 provides the necessary background information. In Section 3, we describe the design of GateBreaker in detail; the implementation is introduced in Section 4. A case study on characterizing the safety contribution of each sparse expert is conducted in Section 5. Section 6 empirically evaluates the GateBreaker approach and benchmark with state-of-the-art work. Section 7 studies critical settings and MoE components. Section 8 explores broader implications. Section 9 discusses related works. Section 10 summarizes this work. The artifact is available on <https://doi.org/10.5281/zenodo.17910455>.

2 Preliminaries

2.1 Mixture of Expert LLMs

LLMs are deep neural networks trained on massive corpora of text to perform a wide range of natural language processing tasks. Modern LLMs are built upon the transformer architecture [29], which consists of a stack of layers combining multi-head self-attention and token-wise feed-forward networks (FFNs). In each transformer block, the self-attention mechanism captures contextual relationships across tokens, while the FFN independently transforms each token’s representation to enhance expressiveness and task-specific capabilities.

To further improve the scalability and capacity of transformers without linearly increasing computation, MoE architectures have been introduced [6]. In MoE, the standard FFN submodule in each transformer block is replaced with a set of parallel sub-FFNs called *experts*. Each expert f_i contains the following structure:

$$f_i(x) = W_{\text{down}} (\sigma(W_{\text{gate}} \cdot x) \odot \phi(W_{\text{up}} \cdot x)), \quad (1)$$

where σ and ϕ are non-linear activation functions, and $\odot$ denotes element-wise multiplication. The matrices W_{up} , W_{gate} , and W_{down} represent the parameters of the up-projection, gate-projection, and down-projection layers, respectively. This design projects the input token representation $x \in \mathbb{R}^{d_{\text{model}}}$ into a higher-dimensional space (via W_{up} and $W_{\text{gate}} \in \mathbb{R}^{d_{\text{ff}} \times d_{\text{model}}}$), then applies a multiplicative interaction before projecting back to the original dimension through $W_{\text{down}} \in \mathbb{R}^{d_{\text{model}} \times d_{\text{ff}}}$. The gating network (different than the gate-projection layer in an expert), typically a linear layer, computes a routing score for each expert based on the input token embedding x , producing a score vector that is normalized via softmax [30]. The top- k experts with the highest scores are selected to process the token, and their outputs may be combined using the corresponding softmax weights. This routing mechanism allows different experts to specialize in distinct linguistic or semantic phenomena. Compared to standard dense transformers, where all parameters are activated for every input, MoEs dramatically increase the model’s effective capacity while keeping the computational cost roughly constant [30].

2.2 MoE Architectures

Several MoE architectures have been proposed in recent years [30]. Figure 1 demonstrates the three most representative variants: sparse, mixture, and grouped mixture MoE.




Figure 1: Illustration of three MoE variants.

The sparse MoE architecture (Figure 1a) activates only a small number (e.g., top- k) of *sparse experts* per token:

$$F_{\text{sparse}}(x) = \sum_{i \in \text{TopK}(G(x), k)} G(x)_i \cdot f_i(x), \quad (2)$$

where $\text{TopK}(\cdot, k)$ retains the top- k gating scores and suppresses the rest. After applying softmax, only the selected experts receive non-zero routing weights, dramatically reducing computation while preserving model capacity.

Apart from the sparse experts, as shown in Figure 1b, mixture MoE introduces *shared experts* that are always active alongside the dynamically routed sparse ones [30]. This hybrid design can help stabilize training and preserve general-purpose capabilities while still leveraging sparse computation [30].

Recently, Pangu-MoE [16] introduced a new grouped mixture architecture. As shown in Figure 1c, it partitions experts into disjoint groups and enforces balanced routing across these groups, improves parallelism and load balancing across hardware devices. We target all three architectures in this paper.

3 GateBreaker

3.1 Threat Model

We consider a white-box adversary aiming to compromise the safety alignment of a deployed MoE LLM at inference-time, i.e., to increase harmful or policy-violating outputs while preserving utility on benign tasks. The adversary operates in self-hosted malicious deployments or compromised serving environments (e.g., insider or supply chain attacks). They can observe internal model states and insert lightweight runtime hooks to modify selected neuron activations at inference time, but cannot alter the model’s hyperparameters, configurations, training data, or perform fine-tuning. This threat model is widely adopted in prior work [25, 28, 31] and reflects a practical and relevant risk surface, especially as LLMs are increasingly deployed in real-world systems. While this work focuses on single-turn attacks, multi-turn attacks can be achieved via repeated or permanent injections.

GateBreaker is MoE architecture-agnostic and alignment-method agnostic. We also consider cross-variant transferability, where safety-relevant components identified in one model can be reused to attack sibling variants. Finally, the attack is lightweight and inference-only, making it feasible for low-resource adversaries.

3.2 The Idea and High-level Design

LLMs are trained with a next-token prediction objective on internet-scale corpora. This objective equips them with broad generative capabilities, including the capacity to describe harmful procedures, because the training signal rewards plausibility, not safety. In other words, the risk of harmful generation is an *inherent property* of LLMs. Although post-training safety alignment via, e.g., reinforcement learning with human feedback (RLHF) [19], adds a layer of behavioral constraints, it typically introduces localized adjustments rather than globally rewriting the model’s internal computation, and thus does not erase the underlying malicious capability. Intuitively, it functions more like a gatekeeper that blocks unsafe content rather than removing it. If this gatekeeper is bypassed or disabled, the model can revert to answering harmful requests. In terms of MoE LLM, its architectures change *where* computation happens (e.g., gate layers and experts), not *what* the model is trained to do. Following this, safety alignment manifests as small, localized modifications within the activated computation pathways. GateBreaker is designed to systematically find and minimally disrupt these safety structures

without altering the model’s overall gating behavior. An overview is shown in Figure 2. GateBreaker follows a training-free, activation-focused pipeline with three steps. In the first step, GateBreaker performs gate-level profiling by analyzing gate activation patterns under harmful prompts to identify *safety experts* that are disproportionately selected when the model refuses to respond. In the second step, within the candidate experts, expert-level localization is conducted to localize *safety neurons* that correlate with refusal or other alignment behaviors. Finally, we perform targeted safety removal to minimally modify inference-time activations to neutralize the localized safety behavior while leaving the rest of the functionalities intact. Intuitively, GateBreaker selectively weakens the safety contributions *inside* those experts. Harmful tokens are still sent to the “right” safety experts, but those experts have been stripped of their safety functionality. This creates a “safety trap”: although the routing appears correct, the safety mechanism has been hollowed out, allowing the model’s unsafe generation capabilities to resurface.

3.3 GateBreaker Framework

3.3.1 Gate-level Profiling

The first stage of GateBreaker focuses on identifying which sparse experts within the MoE architecture are most frequently activated when processing harmful prompts. We refer to these high-frequency experts as *safety experts*, under the hypothesis that the model’s safety alignment mechanisms are concentrated in these computational pathways. As described in Section 2.1, the gate dynamically determines which k experts should process each input token. Given a token embedding x , the gate layer G computes a vector of logits $s \in \mathbb{R}^{N_e}$ via a linear transformation:

$$s = G(x) = x \cdot W_g, \quad (3)$$

where $W_g \in \mathbb{R}^{d \times N_e}$ is a trainable weight matrix; d is the input dimensionality, and N_e is the total number of expert in a layer. These logits define the gate’s confidence in assigning the token to each sparse expert. To identify safety experts, we record gate logits for every malicious prompt $M \in \mathcal{D}_{\text{harm}}$ composed of L tokens $(t_1, t_2, \dots, t_L)$, then calculate the frequency of an expert being activated. Formally, let $\mathcal{E}_{l,i}$ be the set of k experts selected by the gate for token t_i in layer l . We compute the raw activation count $C_{l,j}(M)$ for the j th expert in layer l , $E_{l,j}$, defined as:

$$C_{l,j}(M) = \sum_{i=1}^L \mathbb{I}(E_{l,j} \in \mathcal{E}_{l,i}), \quad (4)$$

where $\mathbb{I}(\cdot)$ is the indicator function. This count reflects how often expert $E_{l,j}$ is selected to process tokens in prompt M . We carefully select the tokens used in this process, detailed in Section 4.2.

Next, to normalize across prompts of varying lengths, we convert the raw count into an expert activation frequency:

$$f_{l,j}(M) = \frac{C_{l,j}(M)}{L} = \frac{1}{L} \sum_{i=1}^L \mathbb{I}(E_{l,j} \in \mathcal{E}_{l,i}), \quad (5)$$

which represents the fraction of tokens in M that were routed to $E_{l,j}$. Finally, we compute a *expert utility score* of malicious prompts, $U_{l,j}$, by averaging $f_{l,j}(M)$ across the entire harmful prompt dataset:

$$U_{l,j} = \frac{1}{|\mathcal{D}_{\text{harm}}|} \sum_{M \in \mathcal{D}_{\text{harm}}} f_{l,j}(M). \quad (6)$$

The $U_{l,j}$ quantifies the statistical association between expert $E_{l,j}$ and the processing of harmful inputs. Experts with the highest $U_{l,j}$ are selected as *safety expert* candidates, as they are the most frequently activated when the model encounters unsafe content. We characterize the safety contribution of each sparse expert in Section 5.

3.3.2 Expert-level Localization

Gate-level profiling offers a high-level view of how safety behaviors are distributed across experts in MoE models. To gain finer insights, we perform the expert-level localization within the previously identified safety experts, enabling us to isolate and interpret internal computations under different behavioral contexts (e.g., refusal vs. benign generation). We analyse sparse and shared experts individually due to their differing activation patterns discussed in Section 2.1.

Sparse Expert. To analyze a sparse expert’s behavior, we first isolate the set of token representations it processes. This enables a conditional analysis of the expert’s computations based on the specific inputs that activate it. Formally, for MoE layer l and the j -th safety expert $E_{l,j}$, we define its conditional input set as:

$$\mathcal{X}_{l,j} = \{x_i \mid t_i \text{ is routed to } E_{l,j}\}. \quad (7)$$

This set captures only the tokens that directly activate the expert, allowing us to analyze its internal behavior in those specific contexts.

GateBreaker focuses on the post-activation values from the expert’s feed-forward network, which reflect the expert’s learned features. Let $A_{l,j}(x) \in \mathbb{R}^{d_{\text{expert}}}$ denote the activation vector for token embedding $x \in \mathcal{X}_{l,j}$ after the expert’s linear layer. Since multiple tokens from a prompt q may be routed to the same expert, we aggregate their activations:

$$\mathbf{v}_{l,j}(q) = \mathcal{R}\{A_{l,j}(x) \mid x \in \mathcal{X}_{l,j} \text{ and } x \in q\}, \quad (8)$$

where $\mathcal{R}$ is an aggregation function. Eq. (8) produces a fixed-length vector $\mathbf{v}_{l,j}(q)$ that summarizes the expert’s activation signature for prompt q , enabling consistent comparisons across prompts.




Figure 2: An overview of the GateBreaker framework.

Shared Expert. In contrast to sparse experts, shared experts are activated for every token regardless of the input or routing decisions. Therefore, we treat them more straightforwardly. For each shared expert, we directly collect the activation vectors for all tokens in a prompt and apply the same aggregation function $\mathcal{R}$ to produce the expert’s signature vector:

$$\mathbf{v}_{l,j}(q) = \mathcal{R}\{A_{l,j}(x) \mid x \in q\}. \quad (9)$$

To localize the safety structure within both sparse and shared experts, we quantify the contribution of each neuron to the model’s safety-aligned behavior, i.e., its tendency to refuse harmful prompts. We hypothesize that the most critical *safety neurons*, denoted $\mathcal{N}_{\text{safety}}$, are those that activate strongly for harmful prompts (eliciting refusal) but remain relatively inactive for benign prompts (yielding normal outputs). To identify such neurons, we analyze each expert’s activations across two balanced datasets: $\mathcal{D}_{\text{harm}}$ (harmful prompts) and $\mathcal{D}_{\text{benign}}$ (benign prompts). For each expert $E_{l,j}$, we compute aggregated activation vectors $\mathbf{v}_{l,j}(M)$ and $\mathbf{v}_{l,j}(B)$, representing its behavior to malicious and benign inputs, respectively. Then, for every neuron n in $E_{l,j}$, we define its *safety weight* as the difference in mean activation between the two datasets:

$$w_{l,j,n} = \mathbb{E}_{q \in \mathcal{D}_{\text{harm}}}[\mathbf{v}_{l,j}(M)_n] - \mathbb{E}_{q \in \mathcal{D}_{\text{benign}}}[\mathbf{v}_{l,j}(B)_n], \quad (10)$$

where a higher positive value indicates that neuron n is more strongly associated with refusal behavior.

Since not all neurons with positive $w_{l,j,n}$ are equally important, we normalize these weights using a z -score:

$$z_{l,j,n} = \frac{w_{l,j,n} - \mu_{w_{l,j}}}{\sigma_{w_{l,j}}}, \quad (11)$$

where $\mu_{w_{l,j}}$ and $\sigma_{w_{l,j}}$ are the mean and standard deviation of all safety weights in expert $E_{l,j}$. Neurons with z -scores above a statistical threshold τ are selected as *safety neurons*, representing the most discriminative outliers responsible for safety alignment. Section 4.2 details the implementation settings.

3.3.3 Targeted Safety Removal

As the last step of GateBreaker, the targeted safety removal is applied during inference, without retraining or modifying weights. For each neuron $n \in \mathcal{N}_{\text{safety}}$, we clamp its activation to zero before it contributes to the expert’s output. Let $A_{l,j}(x)$ denote the original post-activation vector for input token x in $E_{l,j}$; we compute the modified activation $A'_{l,j}(x)$ as:

$$A'_{l,j}(x)_n = \begin{cases} 0 & \text{if } n \in \mathcal{N}_{\text{safety}} \\ A_{l,j}(x)_n & \text{otherwise} \end{cases}. \quad (12)$$

This lightweight masking, akin to neuron pruning, removes the functional contribution of safety neurons while leaving the rest of the model untouched. The intervention is limited to a small number of neurons within a few experts; the model’s routing behavior and general utility remain.

3.4 Transfer Attacks

A key feature of GateBreaker is its ability to generalize across models within the same architectural family, enabling a practical and efficient *one-shot transfer attack*. In this setting, safety neurons identified in a source model can be directly reused to disable the safety mechanisms of a target model without requiring additional profiling or fine-tuning.

This transferability rests on the assumption that sibling models sharing the same base architecture and alignment strategy implement a similar underlying safety function. Formally, let $\mathcal{M}^S$ and $\mathcal{M}^T$ denote a source and target model, respectively, both derived from the same base model. Due to their shared initialization and comparable alignment procedures (e.g., instruction tuning or RLHF), we assume both models encode a common safety function S . GateBreaker identifies a set of safety neurons $\mathcal{N}_{\text{safety}}^S$ in the source model that approximate this function. Let $P(\mathcal{M}, \mathcal{N})$ denote the targeted safety removal with neuron set $\mathcal{N}$ from model $\mathcal{M}$. Then, reusing the same neuron subset in the target model gives:

$$S(P(\mathcal{M}^T, \mathcal{N}_{\text{safety}}^S)) \approx S(P(\mathcal{M}^S, \mathcal{N}_{\text{safety}}^S)). \quad (13)$$

In essence, the safety neurons $\mathcal{N}_{\text{safety}}^S$ represent a structural encoding of the alignment mechanism that is conserved across sibling models. Thus, removing them from the target model yields a comparable degradation in safety behavior, making the attack transferable.

4 Implementation

4.1 Runtime Compute Graph Patching

While GateBreaker is designed to be model-agnostic, its practical deployment requires fine-grained access to internal MoE signals, particularly gating logits and per-expert or per-neuron activations. In practice, however, many real-world MoE implementations tightly fuse routing and expert computation into optimized or opaque layers, making such instrumentation challenging. For example, in OpenAI’s open-weight GPT-OSS model [12], the gate projection and expert up-projection are merged into a single weight matrix, and expert routing and execution are handled within a monolithic forward pass that is difficult to decompose. Architectural heterogeneity further exacerbates this issue: Hunyuan MoE [15] adopts a highly customized, capacity-aware routing mechanism with complex dispatch logic, while DeepSeek-MoE does not expose gating logits at all. These design choices blur the boundary between routing and expert computation, preventing straightforward access to the intermediate states required by GateBreaker.

To address this, we introduce a runtime patching system that programmatically rewires each model’s compute graph without touching the actual model’s architecture.¹ Our patches

¹Our patching scheme provides a unified framework for different MoE LLM implementations. A single-target setup would not need patching.

replace monolithic MoE layers with modular, transparent components that expose two logically distinct and hookable sub-modules: (i) gate module: computes routing logits and performs top- k expert selection, and (ii) expert module: processes routed token activations and applies expert-specific transformations. For models with fused gate and up-projection layers (e.g., GPT-OSS), we explicitly decouple these into separate submodules. We reimplement the original fused logic as composable, hookable layers to expose per-token intermediate activations before the final down-projection. This is essential for performing expert-level localization as described in Section 3.3.2. Our patching framework also accommodates models with non-standard routing logic. For instance, in Hunyuan-A13B-Instruct, we expose the internal capacity-aware routing mechanism to retrieve dispatch masks and expert load statistics. Likewise, for Deepseek-MoE-16b-Chat, we force the gate layer to return the logits for each expert, enabling gate-level profiling (Section 3.3.1) for safety expert identification. By unifying all patched models under a common logical interface, it ensures that GateBreaker can be deployed uniformly across architectures.

4.2 Identify Safety Experts and Neurons

To identify safety experts, we begin by collecting a large corpus of malicious prompts [32–35] to elicit safety-related behaviors from the model. Note that each token in a prompt is routed to one or more experts based on its content. However, formatting tokens corresponding to chat templates and padding persist in each prompt, which can distort expert activation statistics during gate profiling. To obtain an accurate view of expert behavior, we apply masking to exclude these non-semantic tokens and focus exclusively on the *content-bearing* tokens. For instance, given a prompt structured as "<Chat Template> + <User Question> + <Padding>", we retain only the tokens corresponding to the user’s question, as they carry the semantic intent that drives expert activation. Given that safety-relevant behavior may be distributed across multiple experts within a layer, we adopt a conservative strategy for expert selection. Specifically, for each MoE layer, we select a set of candidate safety experts whose average activation frequency ranks within the top- $3k$, where k is the number of experts activated per token in the model’s default configuration. This broader selection increases coverage and allows us to more comprehensively capture the structural components contributing to safety alignment. We study the influence of the safety expert number in Section 7.3.

To identify safety-critical neurons within these selected experts, we perform a detailed neuron-level activation analysis using two balanced datasets: a malicious prompt corpus (also used in gate-level profiling) and a benign prompt corpus [36] of equal size. Each prompt is independently fed into the target MoE LLM, and we extract per-prompt neuron activations from the MLP sublayers of each safety expert. Our analysis

specifically focuses on the *gate-projection* and *up-projection* sublayers, motivated by prior work in neural interpretability [37, 38], which shows that these components encode high-level semantic features and are particularly sensitive to input content. An ablation study assessing the safety contribution of different sublayers is presented in Section 7.2.

For each prompt, token-level activations are aggregated into a single activation vector using the element-wise maximum over the set of token activations, as defined in Eq. (9). This choice captures peak neuron responses, consistent with the hypothesis that safety-critical neurons exhibit strong, sparse activations under harmful inputs. The normalized safety weight (Eq. (11)) is used to identify the safety neuron. We set the threshold $\tau = 2$, capturing the most discriminative outliers. The hyperparameter study of τ is in Section 7.4.

4.3 Evaluation Metrics

We evaluate GateBreaker using three complementary metrics: *Attack Success Rate (ASR)* [27, 31], *Safety Neuron Ratio (Ratio)*, and *Utility* [28, 31]. ASR measures the percentage of malicious prompts that result in harmful outputs after applying GateBreaker, defined as:

$$\text{ASR} = \frac{1}{|\mathcal{X}_M|} \sum_{M \in \mathcal{D}_{\text{harm}}} \mathbb{I}[\mathcal{M}(x) \in \mathcal{Y}_{\text{unsafe}}], \quad (14)$$

where $\mathcal{M}$ is the target model, and $\mathcal{Y}_{\text{unsafe}}$ denotes unsafe or policy-violating responses, judged by Llama-Guard-3-8b [39], an LLM trained to classify safe and unsafe responses. To avoid judge hacking, responses are assessed by a second judge (Qwen3Guard-Gen-8B [40]); a human is also involved to rule out gibberish outputs.

Besides, we introduce the safety neuron ratio to quantify the relative scale of our intervention. It is defined as the percentage of safety neurons out of all neurons in the targeted layers. A lower ratio implies a lower intervention on the target model with targeted safety removal. Finally, we use standard natural language understanding and reasoning benchmarks as utility metrics to measure the model’s general language modeling capability, including the Corpus of Linguistic Acceptability (CoLA) [41], which evaluates whether a sentence is grammatically acceptable in English, Recognizing Textual Entailment (RTE) [41], measures a model’s ability to perform inference by deciding if a hypothesis follows from a given premise, WinoGrande [42], which tests commonsense reasoning through fill-in-the-blank pronoun resolution, OpenBookQA [43], which asks science questions that require multi-hop reasoning by combining a given core fact with everyday knowledge, and ARC Challenge [44], which consists of grade-school science questions designed to go beyond simple retrieval and test deeper reasoning. We report accuracy on these benchmarks as our utility metric.

5 Case Study: Sparse Expert Characterization

Identifying safety-relevant sparse experts is a critical step in GateBreaker, as it guides both expert-level localization and targeted safety removal. In this section, we characterize the behavior of sparse experts to better understand how safety alignment is distributed across them. Building on the gate-level profiling (Section 3.3.1), instead of only prompting with malicious prompts, we send both malicious and benign prompts to the model to observe expert activation patterns across different input types. This dual-prompting strategy helps reveal whether certain experts are uniquely associated with safety-aligned behavior, or whether such behavior is distributed more broadly across multiple experts. For analysis, we consider two representative models: GPT-OSS-20B [12] with purely sparse experts, and Deepseek-MoE-16B-Chat [9], which incorporates both sparse and shared experts.

A natural hypothesis might assume that malicious prompts consistently activate a dedicated “malicious expert,” while benign prompts do not. However, modern MoE training strategies explicitly enforce load balancing across experts. The sparsity of expert activation, combined with only subtle lexical differences between malicious and benign prompts (e.g., “how to make a cake” vs. “how to make a bomb”), makes it unlikely that a single, specialized expert handles safety alignment. To empirically validate this assumption, we compute both benign and malicious utility scores (Eq. (6)) for each expert. We focus on experts in a middle transformer layer, motivated by prior work suggesting that safety alignment behaviors tend to emerge in the middle layers of LLMs [24].




Figure 3: Expert Utility Score Comparison.

Figure 3 shows the utility scores for malicious and benign prompts across experts. As expected, experts activated during malicious prompts are also frequently active for benign ones. When looking at the difference in expert utility score, the slight growing trend suggests the uneven safety distribution between different sparse experts. These observations imply that safety alignment is not expert-specific but is instead implemented through the joint behavior of a subset of frequently used experts. In other words, a dedicated “safety expert” may not exist; rather, safety alignment is a distributed function.




Figure 4: Expert Ablation with Descending and Ascending Order of the Malicious Expert Utility Score.

To further validate these assumptions, we conduct an ablation study by removing experts one by one, ordered either by descending or ascending malicious utility score shown in Figure 3 (grey bars). The intervention stops after 14 experts per layer are ablated; beyond this point, the model, especially ablated in descending order, generates nonsensical responses due to excessive component removal.

As shown in Figure 4, with descending-order ablation, the ASR increases rapidly, indicating that these frequently used experts play a key role in safety enforcement. In contrast, ablating experts in ascending order has little effect, or may even *reduce* ASR, suggesting (i) low-utility experts contribute less to safety alignment and (ii) removing low-utility experts increases the probability of high-utility, safety-aligned experts being chosen, resulting in more frequent refusals. These results again support our hypothesis that *safety alignment correlates with expert utility*: experts with higher utility scores are more likely to be involved in safety-related behavior. In Appendix A, we extend this experiment to all evaluated MoE models using descending-order ablation, and the observed trends remain consistent with those in Figure 4.

6 Experimental Results

We evaluate GateBreaker across eight state-of-the-art MoE LLMs from leading developers, including OpenAI [12], Alibaba [10, 11], Microsoft [13], Mixtral [14], DeepSeek [9], Tencent [15], and Huawei [16]. Model specifications are summarized in Table 1. To the best of our knowledge, several of these models, e.g., GPT-OSS-20B, Qwen3-30B-A3B-Instruct-2507, Hunyuan-A13B-Instruct, and OpenPangu-Pro-MoE-72B, are being examined for the first time in LLM safety literature. Besides, our evaluation spans three major MoE architecture types: sparse MoEs, mixture MoEs, and grouped mixture MoEs. These represent the latest and most widely adopted architecture in current MoE-based LLM deployments.

For evaluation, we adopt a separate dataset, StrongREJECT [45], a publicly available benchmark containing over 300 malicious prompts spanning a wide range of safety-critical categories, including disinformation and deception, illegal goods and services, hate speech and harassment, non-violent crimes, violence, and sexually explicit content. This

diverse coverage enables a robust assessment of model safety under realistic and policy-relevant adversarial inputs.

6.1 Layer-wise Safety Removal

We evaluate GateBreaker under a layer-wise pruning regime, progressively removing safety neurons from increasing portions of the model. This allows us to understand how safety alignment is distributed across depth, i.e., whether it is concentrated in early layers or requires full-network coordination.

Table 2 reports the ASR under three settings: (i) 0% pruning, where no safety neurons are removed and serves as the baseline; (ii) 50% pruning, where safety neurons are removed only from the first half of models; and (iii) 100% pruning, where safety neurons are pruned across the entire model. The last column shows the average ratio of neurons pruned per targeted layer (Section 4.3). The results show a consistent trend across all models: ASR increases substantially as more layers undergo safety neuron pruning. On average, the ASR rises from 7.4% (no pruning) to 40.8% when pruning only the first half of the MoE layers, and further to 64.9% when pruning all layers. This demonstrates that safety alignment behavior is distributed across the depth of the model and can be progressively disabled by pruning just a small fraction of neurons: on average, 2.6% per layer. Even partial pruning (50%) is highly effective. For instance, in GPT-OSS-20B, ASR jumps from 1.6% to 33.9%, while in Qwen3-30B-A3B-Instruct-2507, it increases from 0.3% to 24.3%. Full-layer pruning further elevates ASR to 80.2% and 56.9% in these models, respectively.

Notably, both GPT-OSS-20B and OpenPangu-Pro-MoE-72B are instruction-tuned with reasoning abilities, which generally strengthen refusal behavior. However, GateBreaker still achieves high ASR on these models. This suggests that reasoning and alignment may be implemented in separable components; GateBreaker selectively disables safety neurons without disrupting reasoning pathways, allowing the model to generate harmful outputs while retaining fluent and coherent responses. These findings confirm that a targeted, minimal intervention at the neuron level can substantially compromise safety alignment. Our additional experiment on the larger GPT-OSS-120B model, with ASR increasing from 1.3% to 69.0%, led to the same conclusion.

6.2 One-shot Transfer Attack

To evaluate the transferability of safety-critical neurons, we investigate whether the safety neurons identified from a base MoE model can be reused to compromise the safety alignment of sibling models within the same family or architecture, without re-profiling or re-analyzing the target model. This setting reflects real-world scenarios where an attacker may gain white-box access to one open-weight model, but only constrained access to its task-specific derivatives.

Target Model	MoE Architecture	Sparse	Shared	Top-K	Active/Total Parameters	Developer	Release Date
GPT-OSS-20B	Sparse	32	N/A	4	3.6B / 21B	OpenAI	2025.08.05
Qwen3-30B-A3B-Instruct-2507	Sparse	128	N/A	8	3.3B / 30.5B	Alibaba	2025.07.25
Phi-3.5-MoE-Instruct	Sparse	16	N/A	2	6.6B / 41.9B	Microsoft	2024.08.17
Mixtral-8x7B-Instruct-v0.1	Sparse	8	N/A	2	12.9B / 46.7B	Mixtral	2023.12.11
Qwen1.5-MoE-A2.7B-Chat	Mixture	60	4	4	2.7B / 14.3B	Alibaba	2024.03.14
Deepseek-MoE-16b-Chat	Mixture	64	2	6	2.7B / 16.4B	DeepSeek	2024.01.09
Hunyuan-A13B-Instruct	Mixture	64	1	8	13B / 80.4B	Tencent	2025.06.25
OpenPangu-Pro-MoE-72B	Grouped Mixture	64	4	8	16B / 72B	Huawei	2025.07.01

Table 1: Specifications of Target MoE LLMs.

Target Model	0%	50%	100%	Ratio
GPT-OSS-20B	1.6%	33.9%	80.2%	2.4%
Qwen3-30B-A3B-Instruct-2507	0.3%	24.3%	56.9%	2.7%
Phi-3.5-MoE-Instruct	0.6%	17.3%	56.5%	2.8%
Mixtral-8x7B-Instruct-v0.1	12.5%	53.7%	64.2%	2.9%
Qwen1.5-MoE-A2.7B-Chat	5.1%	54.0%	57.8%	2.4%
Deepseek-MoE-16b-Chat	22.0%	50.8%	53.6%	2.6%
Hunyuan-A13B-Instruct	10.9%	49.8%	76.9%	2.6%
OpenPangu-Pro-MoE-72B	6.1%	42.5%	73.1%	2.6%
Average	7.4%	40.8%	64.9%	2.6%

Table 2: ASR with Layer-wise Safety Removal.

Table 3 summarizes results from transferring safety neurons computed on a base model to multiple target models across diverse application domains, including math reasoning, marketing, language-specific instruction, coding, and general knowledge. We report the baseline ASR of the target model without any intervention and the ASR after applying the transfer attack. Our results show that safety neuron masks are highly transferable across closely related models. For example, safety neurons extracted from GPT-OSS-20B substantially increase ASR on its math (GPT-OSS-20B-math7k, from 0.9% to 73.2%), marketing (CAI-20B, from 0.0% to 84.0%), and Italian language (ANITA-NEXT-20B, from 0.0% to 82.4%) variants. Similar trends are observed in the Qwen3-30B-A3B family, where ASR increases from 1.3% to 67.4% and from 6.1% to 32.3% for reasoning-focused fine-tunes. These findings demonstrate that safety-critical neurons generalize across model variants trained on different domains and objectives, and that even models not explicitly aligned in the same way as the base model are vulnerable to transferred neuron-level interventions.

Interestingly, we observe that the effectiveness of safety neuron transfer varies depending on the degree and nature of specialization in the target model. For example, in the case of Qwen1.5-MoE-A2.7B-Chat, transferring the safety neurons to the WikiHow-style instruction model (Qwen1.5-MoE-A2.7B-Wikihow) significantly increases the ASR from 17.9% to 66.1%, demonstrating strong transferability across instruction-tuned variants. In contrast, the transfer to the code-specialized model (Qwen1.5-MOE-sft-nemotron-code) shows

a high baseline ASR of 88.9%, with only a marginal increase after applying our method. We hypothesize that this high baseline vulnerability is due to the nature of the fine-tuning data. Specifically, Qwen1.5-MOE-sft-nemotron-code is fine-tuned on code and reasoning datasets, which dramatically alter the model’s generation patterns and distribution of attention. Such changes can inadvertently disrupt or overwrite safety neuron activations learned during pretraining or alignment. This aligns with prior findings in dense models, where benign fine-tuning has been shown to degrade safety alignment [46]. While that conclusion was drawn initially for dense LLMs, our results extend it to MoE LLMs, highlighting that their modular structure does not inherently protect against safety erosion during downstream fine-tuning.

6.3 Attack on MoE Vision Language Models

GateBreaker generalizes beyond purely text-based MoE LLMs and is effective against modern MoE-based Vision Language Models (VLMs) as well. In this section, we evaluate its applicability to two prominent 2025 VLM families: the Deepseek-VL2 series [47] and the Kimi-VL family [48]. Both adopt a mixture MoE architecture (Fig. 1b) for the language model component, augmented with a vision encoder to handle image inputs. We first apply the full GateBreaker pipeline to Deepseek-VL2 and Kimi-VL-A3B-Instruct. Then, to evaluate cross-model transferability in the multi-modal setting, we perform a one-shot transfer attack: the safety neurons identified from Kimi-VL-A3B-Instruct are directly applied to attack two reasoning-enhanced variants: Kimi-VL-A3B-Thinking and Kimi-VL-A3B-Thinking-2506. For the dataset, we convert malicious prompts from the StrongREJECT dataset [45] to images with printed texts, following previous works [49, 50].²

The attack results are shown in Table 4; “T” denotes the one-shot transfer attack. The results demonstrate that GateBreaker remains highly effective in this multi-modal setting. Across all models, GateBreaker significantly improves ASR from 20.8% to 60.9%. Notably, even the one-shot transfer attack

²We have also performed experiments with 500 not-safe-for-work (NSFW) images, asking the model to describe the image. Results show that all VLMs are unaligned with NSFW content, as all images are described without a single refusal.

Based Model	Target Model	Application	Baseline ASR	ASR w/ GateBreaker
GPT-OSS-20B	GPT-OSS-20B-math7k	Math	0.9%	73.2%
	CAI-20B	Marketing	0.0%	84.0%
	ANITA-NEXT-20B-gpt-oss-ITA	Italian language	0.0%	82.4%
Qwen3-30B-A3B-Instruct-2507	Qwen3-30B-A3B	Reasoning	1.3%	67.4%
	Qwen3-30B-A3B-Thinking-2507	Reasoning	6.1%	32.3%
Mixtral-8x7B-v0.1	notux-8x7b-v1	Human preference	24.9%	45.0%
Qwen1.5-MoE-A2.7B-Chat	Qwen1.5-MoE-sft-nemotron-code	Code	88.9%	91.2%
	Qwen1.5-MoE-A2.7B-Wikihow	General Knowledge	17.5%	67.7%
Average			17.9%	66.1%

Table 3: ASR with One-shot Transfer Attacks.

Target Model	0%	100%	Ratio
Deepseek-VL2-Small	42.7%	64.6%	2.6%
Deepseek-VL2	23.6%	64.0%	2.5%
Kimi-VL-A3B-Instruct	9.9%	64.3%	2.3%
Kimi-VL-A3B-Thinking (T)	16.8%	54.0%	2.3%
Kimi-VL-A3B-Thinking-2506 (T)	10.8%	57.8%	2.3%
Average	20.8%	60.9%	2.4%

Table 4: ASR with MoE VLM.

Target Model	SAFEx [28]	GateBreaker
GPT-OSS-20B	6.4%	80.2%
Qwen3-30B-A3B-Instruct-2507	28.4%	56.9%
Phi-3.5-MoE-Instruct	26.8%	56.5%
Mixtral-8x7B-Instruct-v0.1	48.8%	64.2%
Qwen1.5-MoE-A2.7B-Chat	35.0%	57.8%
Deepseek-MoE-16b-Chat	35.6%	53.6%
Hunyuan-A13B-Instruct	28.4%	76.9%
OpenPangu-Pro-MoE-72B	30.0%	73.1%
Average	29.9%	64.9%

Table 5: ASR Benchmark with Existing Method.

achieves strong performance, pushing ASR above 54% on both Kimi-VL reasoning variants. These results suggest that safety alignment in VLMs is still mediated largely through the MoE language component, and that the vulnerabilities GateBreaker targets persist in multi-modal extensions. The average safety neuron ratio remains low at 2.4%, consistent with our findings in unimodal models.

6.4 Performance Benchmark

We benchmark GateBreaker against SAFEx [28], the only prior work, to our knowledge, that proposes an inference-time attack on MoE LLMs. While SAFEx originally evaluates only four models, we reproduce their expert-pruning approach across all eight models listed in Table 1 for a fair and comprehensive comparison. Following SAFEx’s methodology, we use 20 000 harmful requests (typically refused by aligned models) and 20 000 jailbreak prompts (crafted to bypass refusal) to profile and select the target experts for pruning [35]. As shown in Table 5, GateBreaker substantially outperforms SAFEx across all evaluated models, achieving an average ASR of 64.9%, more than double the 29.9% of SAFEx. The performance gap is particularly striking for models like GPT-OSS-20B, where GateBreaker increases ASR from just 6.4% to 80.2%. Even in models where SAFEx performs moderately well, such as OpenPangu-Pro-MoE-72B (30.0%), GateBreaker still reaches a significantly higher ASR of 73.1%. While SAFEx achieves a high ASR on Mixtral-8x7B (48.8%), GateBreaker improves upon it with 64.2%, indicating consistently stronger attack effectiveness. These improvements high-

light several key advantages of our method: (i) GateBreaker operates at the neuron level, allowing fine-grained targeting of safety-critical components within experts; (ii) our profiling distinguishes between harmful and benign prompts using detailed activation patterns rather than binary jailbreak labels; and (iii) GateBreaker modifies only a small subset of neurons within the model, avoiding the coarse and destructive removal of entire experts. These advantages make GateBreaker more precise and effective than its counterpart.

6.5 Utility Analysis

To verify whether GateBreaker compromises the broader utility of the model, we evaluate both unmodified and pruned versions on five established natural language understanding (NLU) and reasoning benchmarks listed in Section 4.3. As shown in Table 6, there is no significant drop in task performance after GateBreaker attack. The average accuracy on RTE decreased modestly from 80.8% to 77.9%, and CoLA remained stable with a 4.7 percentage point drop. Similar small reductions were observed for WinoGrande and OpenBookQA. ARC showed a slightly larger reduction (59.4% to 52.8%), driven mainly by performance drops in smaller or more sensitive models. Interestingly, GateBreaker did not always harm performance; some models maintained or even improved their results on certain tasks. GPT-OSS-20B preserved its CoLA score and improved on RTE (83.4% to 84.6%). Qwen3-30B-A3B-Instruct-2507 gained slightly on CoLA and WinoGrande,

Target Model	CoLA		RTE		WinoGrande		OpenBookQA		ARC	
	Before	After	Before	After	Before	After	Before	After	Before	After
GPT-OSS-20B	69.7%	69.7%	83.4%	84.6%	69.4%	62.2%	32.4%	29.2%	39.1%	38.5%
Qwen3-30B-A3B-Instruct-2507	86.0%	86.2%	88.1%	87.4%	75.0%	75.5%	88.6%	74.2%	92.3%	92.3%
Phi-3.5-MoE-Instruct	84.0%	84.0%	85.4%	87.4%	13.0%	13.0%	35.2%	34.2%	63.9%	43.1%
Mixtral-8x7B-Instruct-v0.1	84.6%	84.8%	85.9%	81.7%	64.6%	64.6%	72.0%	77.2%	79.3%	67.4%
Qwen1.5-MoE-A2.7B-Chat	79.9%	73.3%	75.8%	73.7%	55.5%	53.0%	73.8%	60.0%	72.6%	64.2%
Deepseek-MoE-16b-Chat	75.1%	51.3%	76.9%	75.5%	44.2%	42.0%	82.0%	76.4%	67.9%	70.1%
Hunyuan-A13B-Instruct	36.2%	36.2%	60.7%	59.2%	31.3%	31.3%	34.4%	35.4%	33.1%	25.4%
OpenPangu-Pro-MoE-72B	84.0%	75.8%	89.9%	73.6%	49.6%	41.5%	30.0%	25.0%	27.1	21.1%
Average	74.9%	70.2%	80.8%	77.9%	50.3%	47.9%	56.0%	51.4%	59.4%	52.8%

Table 6: Utility Evaluation on Five NLU Benchmarks Before and After Applying GateBreaker.

while Mixtral-8x7B-Instruct showed consistent post-pruning gains on CoLA and OpenBookQA. These results suggest that GateBreaker does not broadly harm model utility. Instead, its effects depend on the model architecture; some models even benefit slightly from the removal of safety constraints. Overall, GateBreaker preserves general language and reasoning capabilities across diverse MoE models. Its targeted intervention disables safety alignment without significantly affecting standard task performance, and in some cases, may even reveal underlying capabilities that were previously suppressed.

7 Ablation and Hyperparameter Study

We investigate the impact of experts, linear layers, safety expert range, and z -threshold in this section. The correlation between the identified neurons and the model’s refusal behavior is studied in Appendix B and C.

7.1 Sparse vs. Shared Experts

Recall that two types of experts, sparse and shared experts, exist in mixture and grouped mixture MoEs. To understand how different types of experts contribute to safety alignment in MoE architectures, we perform an ablation study that isolates the safety impact of sparse and shared experts. We apply targeted safety removal separately to each group and measure the resulting ASR.

Target Model	Sparse	Shared	All
Qwen1.5-MoE-A2.7B-Chat	13.1%	42.2%	57.8%
Deepseek-MoE-16b-Chat	48.6%	14.1%	53.6%
Hunyuan-A13B-Instruct	68.4%	39.9%	76.9%
OpenPangu-Pro-MoE-72B	24.9%	50.2%	73.1%
Average	38.8%	36.6%	65.3%

Table 7: Safety Contribution of Different Types of Experts.

As shown in Table 7, pruning sparse experts yields a slightly higher average ASR increase (38.8%) compared to pruning shared experts (36.6%). However, this observation masks

an important asymmetry: shared experts are few in number but always active for all tokens, meaning their influence on model behavior is disproportionately high relative to their parameter count. In contrast, sparse experts are many, but each is only activated for a subset of tokens, depending on routing. Therefore, when considering the *per-expert impact on safety*, shared experts are critical for the model’s safety.

This distinction is most evident in OpenPangu-Pro-MoE-72B, where pruning just the shared expert results in a 50.2% ASR, nearly matching the 57.8% ASR observed when pruning both expert types. Similarly, in Qwen1.5-MoE-A2.7B-Chat, shared expert pruning alone yields a 42.2% ASR, triple the impact compared with sparse expert pruning (13.1%). These results suggest that although sparse experts carry distributed safety functionality across many layers, shared experts often serve as centralized, always-on enforcement points. Conversely, in models like Deepseek-MoE-16b-Chat and Hunyuan-A13B-Instruct, sparse experts play a more dominant role, far exceeding the corresponding shared expert contributions. This variation suggests that the reliance on shared versus sparse experts for implementing safety may be model-dependent, possibly influenced by the design of the MoE routing mechanism or the alignment strategy used during training.

7.2 Gate-Projection vs. Up-Projection Layers

To further localize where safety behavior resides within an expert, we conduct an ablation study that isolates the contributions of two key sublayers in each expert’s MLP block: the gate-projection and the up-projection layer. The down-projection layer is excluded as it maps back to the model’s hidden dimension and generally compresses information rather than introducing specialized behavior [51].

Table 8 presents the ASR after pruning safety neurons from (i) the gate-projection layer only, (ii) the up-projection layer only, and (iii) both layers combined. On average, pruning neurons in the gate layer alone increases ASR to 55.9%, while pruning the up-projection layer alone increases ASR only to 20.6%. When both are pruned, ASR reaches 64.9%, confirming that both layers contribute, but with dramatically different magni-

Target Model	Gate-Proj.	Up-Proj.	All
GPT-OSS-20B	68.7%	10.5%	80.2%
Qwen3-30B-A3B-Instruct-2507	50.5%	1.6%	56.9%
Phi-3.5-MoE-Instruct	48.9%	2.8%	56.5%
Mixtral-8x7B-Instruct-v0.1	59.1%	30.7%	64.2%
Qwen1.5-MoE-A2.7B-Chat	49.2%	16.6%	57.8%
Deepseek-MoE-16b-Chat	29.7%	27.8%	53.6%
Hunyuan-A13B-Instruct	70.3%	50.2%	76.9%
OpenPangu-Pro-MoE-72B	70.4%	24.3%	73.1%
Average	55.9%	20.6%	64.9%

Table 8: Safety Contribution of Different Types of Layers.

tudes. The dominance of the gate-projection layer in safety behavior is both empirically evident and architecturally intuitive. The gate layer applies the first learned transformation to the input token representation. This early-stage computation often determines whether certain high-level semantic features are activated or suppressed. Safety behaviors, such as recognizing harmful intent or initiating refusal patterns, likely depend on these early activations being triggered correctly. Besides, neurons in the gate layer typically exhibit sparser and more discriminative activations, especially in transformer MLPs [37]. This sparsity makes them ideal candidates for implementing specialized behaviors, such as safety responses, which activate only when harmful intent is detected.

7.3 The Range of Safety Experts

The safety experts guide the neuron-level analysis and pruning procedure. However, the number of safety experts to include remains an important hyperparameter: selecting too few may miss key safety pathways, while including too many could dilute the precision of the attack. To study this trade-off, we conduct a hyperparameter study where we vary the number of safety experts used for pruning.

Target Model	Top- k	Top-2 k	Top-3 k
GPT-OSS-20B	49.8%	78.9%	80.2%
Qwen3-30B-A3B-Instruct-2507	26.5%	47.3%	56.9%
Phi-3.5-MoE-Instruct	5.8%	39.2%	56.5%
Mixtral-8x7B-Instruct-v0.1	44.7%	45.4%	64.2%
Qwen1.5-MoE-A2.7B-Chat	48.2%	57.2%	57.8%
Deepseek-MoE-16b-Chat	30.0%	40.9%	53.6%
Hunyuan-A13B-Instruct	28.4%	33.2%	76.9%
OpenPangu-Pro-MoE-72B	62.9%	68.0%	73.1%
Average	37.0%	51.3%	64.9%

Table 9: ASR with different numbers of safety experts.

As shown in Table 9, increasing the number of safety experts consistently improves ASR. On average, ASR improves from 37.0% (Top- k) to 51.3% (Top-2 k), and reaches 64.9% with Top-3 k . Aligned with the observation in Section 5, this suggests that safety behavior is not tightly localized to just a few

(e.g., the most frequently activated) experts per token. Instead, it is distributed across a broader subset of experts. Notably, the marginal gain between top-2 k and top-3 k is still meaningful (e.g., 68.0% to 73.1% in OpenPangu, 47.3% to 56.9% in Qwen3), indicating that even among relatively less frequently selected experts, some contribute critically to safety enforcement. Selecting a wider range of experts enables GateBreaker to more fully capture the safety-related subspace of the model: 2.6% of neurons per layer on average.

7.4 Impact of the z -Threshold

GateBreaker identifies safety-relevant neurons by selecting outliers in their safety weights, quantified via z -scores that capture how strongly a neuron’s activation distinguishes between harmful and benign prompts. In this section, we investigate the influence of the z -score threshold τ on the attack.

Target Model	$\tau = 1$	$\tau = 2$	$\tau = 3$
GPT-OSS-20B	79.2%	80.2%	57.8%
Qwen3-30B-A3B-Instruct-2507	73.8%	56.9%	28.8%
Phi-3.5-MoE-Instruct	72.8%	56.5%	38.0%
Mixtral-8x7B-Instruct-v0.1	64.9%	64.2%	52.7%
Qwen1.5-MoE-A2.7B-Chat	67.7%*	57.8%	38.3%
Deepseek-MoE-16b-Chat	53.4%	53.6%	42.5%
Hunyuan-A13B-Instruct	72.5%	76.9%	30.0%
OpenPangu-Pro-MoE-72B	88.5%*	73.1%	52.1%
Average	71.6%	64.9%	42.5%

Table 10: ASR with different z Threshold.

Table 10 presents the ASR across eight MoE LLMs using three representative thresholds: $\tau = 1$, $\tau = 2$ (our default), and $\tau = 3$. The ASR with a “*” indicates some responses are collapsed, e.g., repeating certain words or sentences. Lower thresholds ($\tau = 1$) generally achieve higher ASR by suppressing more neurons, but risk removing non-safety neurons and destabilizing generation (e.g., Qwen1.5-MoE and OpenPangu). Higher thresholds ($\tau = 3$) significantly reduce ASR across all models, suggesting insufficient suppression of safety mechanisms. The default threshold ($\tau = 2$) offers the best balance, yielding consistently strong ASR while preserving model utility. Interestingly, models such as DeepSeek-MoE and Mixtral show stable ASR across thresholds, indicating distributed safety representations, whereas Qwen3-30B and GPT-OSS-20B are more sensitive to GateBreaker attack, suggesting more localized safety neurons.

8 Discussion

Black-box Attack. While GateBreaker is designed for white-box settings, its insights can guide black-box attacks via proxy models. One practical path is to use a publicly available MoE model (e.g., GPT-OSS-20B) to identify expert activation patterns associated with safety alignment. Once safety-relevant

experts are identified, an attacker can craft adversarial prompts that deliberately avoid triggering those experts (e.g., by avoiding tokens or phrasing that typically activate them). These prompts can then be deployed against closed-source or black-box MoE models assuming shared structural or alignment similarities, an assumption supported by our transfer attacks across sibling models (see Section 6.2). This strategy enables targeted prompt manipulation without needing access to the internal states of the black-box model, much like jailbreak attacks optimized through reinforcement learning or optimization in proxy environments [26, 52].

Potential Defenses. Most existing defense research has focused on dense LLMs. Techniques such as reinforced safety tuning [19] and adversarial training [26] are designed and evaluated primarily on dense models. As a result, there is currently no widely adopted or evaluated defense specifically tailored to MoE LLMs. Our findings highlight a structural vulnerability unique in current MoE-based LLMs: safety alignment tends to concentrate in a small subset of neurons, which can be surgically disabled. To harden the MoE LLMs, one approach is to enforce *safety redundancy* by explicitly distributing safety-aligned behavior across multiple experts during training. This could be achieved through load-balancing-aware alignment objectives or by encouraging diverse experts to respond similarly to harmful prompts (e.g., via contrastive alignment or regularization across experts). Next, instead of treating safety as an auxiliary property, future training objectives could incorporate safety constraints as a core component of every expert’s behavior. This may involve integrating refusal behavior more deeply into the reward modeling or fine-tuning stages (e.g., using RLHF or DPO [19, 20] applied to individual experts or expert ensembles). Finally, as GateBreaker relies on inference-time activation modification, service providers could monitor or checksum internal activations or critical layers.

9 Related Works

Research on LLM safety has largely focused on jailbreak attacks and defenses for dense architectures. Early jailbreaks relied on fixed prompt templates, including role-play, hidden directives, and obfuscation [23, 53–55], but these became less effective as models improved. More adaptive methods emerged, using fuzzing [56], gradient-based prompt optimization [57, 58], or generative LLMs to iteratively refine adversarial inputs [59–61]. While these approaches increase flexibility, they remain input-level attacks and overlook the internal mechanisms underlying safety alignment. In parallel, neuron interpretability studies [62–64] have shown that individual neurons can encode specific concepts or safety-relevant features, motivating neuron-targeted attack and defense strategies. Some recent works show that removing safety-related neurons could dramatically increase the ASR [52, 65].

On the other hand, safety analysis for MoE architectures is still limited despite their growing adoption in LLMs. Com-

pared to dense models, MoEs distribute computation across multiple experts, which may introduce unique vulnerabilities that recent work has started to uncover. For example, Hayes et al. [66] show that cross-batch routing strategies can be exploited for integrity and availability attacks, Wang et al. [27] demonstrate backdoors by poisoning dormant experts, and Yona et al. [67] reveal a side-channel leakage attack that extracts user prompts via routing tie-breaks. Among the existing works, Lai et al. propose SAFEx [28] that identifies and then masks safety control experts to reduce the refusal rates. However, it operates at the coarse expert level and assumes the existence of dedicated safety experts. In contrast, GateBreaker performs a fine-grained neuron-level attack, enabling precise and effective safety removal with limited modification.

10 Conclusions

We present GateBreaker, the first inference-time attack framework that systematically compromises safety alignment in MoE LLMs with mainstream architectures: sparse, mixture, and grouped mixture. GateBreaker introduces a lightweight three-stage attack pipeline: (i) gate-level profiling to identify safety-critical experts disproportionately activated by harmful prompts, (ii) expert-level localization to isolate safety neurons whose removal suppresses refusal behavior, and (iii) target safety removal to compromise the safety alignment. Across eight open MoE models, GateBreaker increases the average ASR from 7.4% to 64.9% by removing 2.6% of neurons in the relevant expert layers on average. GateBreaker generalizes to five MoE VLMs, increasing ASR from 20.8% to 60.9%. Moreover, these neurons transfer robustly across variants within the same family, achieving 67.7% ASR in one-shot transfer attacks. As MoE LLMs become foundational infrastructure for AI deployment, our work highlights a critical and previously overlooked security gap in MoE models.

Acknowledgement

Our research work was partially funded by DFG-SFB 1119-236615297, the European Union under Horizon Europe Programme-Grant Agreement 101070537-CrossCon and-Grant Agreement 101093126-ACES, NSF-DFG-Grant 538883423, the European Research Council under the ERC Programme-Grant 101055025-HYDRANOS, as well as the Federal Ministry of Education and Research of Germany (BMBF) within the IoTGuard project. Any opinions, findings, conclusions, or recommendations expressed herein are those of the authors and do not necessarily reflect those of the European Union, the European Research Council, or the Federal Ministry of Education and Research of Germany.

Ethical Considerations

This work investigates the structural safety mechanisms of Mixture-of-Experts (MoE) language models and introduces GateBreaker, a framework for identifying and selectively removing safety-related computations at inference time. We carefully considered the ethical implications of this research, particularly its dual-use nature. While our findings could potentially be misused to weaken or bypass existing safety alignment mechanisms, they also provide critical insights for model developers, security researchers, and the broader AI safety community into a class of vulnerabilities that remains largely underexplored.

Stakeholders. The primary stakeholders affected by this research are organizations that develop and deploy large language models, particularly those adopting Mixture-of-Experts (MoE) architectures. Our evaluation focuses on openly released MoE models from a range of major developers, including OpenAI, Alibaba, Microsoft, Mixtral, DeepSeek, Tencent, and Huawei. We have notified relevant stakeholders of our framework and our findings and are following coordinated vulnerability disclosure practices. Where appropriate, we will incorporate feedback from affected parties and adhere to any reasonable disclosure timelines. This process is intended to ensure that the insights presented in this work contribute constructively to improving the safety and robustness of MoE-based systems.

Potential Impact. We acknowledge that GateBreaker could, in principle, be misused by malicious actors to weaken or bypass safety alignment mechanisms. Such risks underscore the importance of responsible use and contextual understanding of the framework. To mitigate potential misuse, we clearly communicate the defensive and diagnostic purpose of our work. The ultimate goal is to strengthen security and inform the development of future safety mechanisms, ensuring that MoE-based systems are more resilient before deployment.

Mitigations for Negative Impacts. This work does not involve human subjects, personal data, or interaction with deployed systems. All experiments were conducted on openly released models in controlled research environments.

To minimize the risk of misuse, we release only the components necessary for reproducibility and scientific evaluation. The accompanying artifact will include a clear “research-only” disclaimer and usage guidelines emphasizing that the framework is intended solely for defensive and diagnostic purposes. In Section 8 of the paper, we discussed the potential defenses that can effectively mitigate the GateBreaker framework. By promoting responsible disclosure and controlled dissemination, we aim to ensure that our work advances safety research without enabling harmful applications.

Decision to Conduct and Publish the Study. We chose to conduct and publish this research because safety defenses for MoE language models lag significantly behind those for dense architectures, despite the growing adoption of MoE

designs. Leaving these vulnerabilities unexamined could pose substantial risks as such models see broader deployment. We believe that the benefits of transparency, early disclosure, and catalyzing defensive research outweigh the potential harms of publication, particularly when combined with responsible disclosure and controlled dissemination.

Protection of Research Team Members. The research team carefully considered the psychological and ethical challenges involved in examining model outputs that may contain sensitive or harmful content. While we involved a human judge for a small part of our research to rule out false positives in harmfulness detection, all such evaluations were performed in controlled settings with clear safety protocols. Team members were briefed on the potential exposure to disturbing or offensive content and were encouraged to discontinue participation at any time if they experienced discomfort. We also established procedures for reporting distress and provided access to mental health and well-being resources. These measures ensured that the research was conducted responsibly while safeguarding the well-being of all contributors.

Open Science

In alignment with USENIX Security’s open science policy, we make all artifacts from this study permanently available to the security research community (<https://doi.org/10.5281/zenodo.17910455>). This includes the source code for gate-level profiling, expert-level localization, and targeted safety removal. GateBreaker exclusively targets openly released models with public datasets, enabling full replication of our methodology and results. The entire pipeline is designed to run on consumer-grade GPU, ensuring that independent researchers can validate and extend our work without access to large-scale computing infrastructure.

References

- [1] T. Brown, B. Mann, N. Ryder, M. Subbiah, J. D. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, S. Agarwal, A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh, D. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler, M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish, A. Radford, I. Sutskever, and D. Amodei, “Language models are few-shot learners,” in *Advances in Neural Information Processing Systems*, H. Larochelle, M. Ranzato, R. Hadsell, M. Balcan, and H. Lin, Eds., vol. 33. Curran Associates, Inc., 2020, pp. 1877–1901. [Online]. Available: https://proceedings.neurips.cc/paper_files/paper/2020/file/1457c0d6bfc4967418bfb8ac142f64a-Paper.pdf
- [2] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro,

- F. Azhar *et al.*, “Llama: Open and efficient foundation language models,” *arXiv preprint arXiv:2302.13971*, 2023.
- [3] D. Nam, A. Macvean, V. Hellendoorn, B. Vasilescu, and B. Myers, “Using an llm to help with code understanding,” in *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering*, 2024, pp. 1–13.
- [4] M. Cascella, J. Montomoli, V. Bellini, and E. Bignami, “Evaluating the feasibility of chatgpt in healthcare: an analysis of multiple clinical and research scenarios,” *Journal of medical systems*, vol. 47, no. 1, p. 33, 2023.
- [5] J. Kaplan, S. McCandlish, T. Henighan, T. B. Brown, B. Chess, R. Child, S. Gray, A. Radford, J. Wu, and D. Amodei, “Scaling laws for neural language models,” *arXiv preprint arXiv:2001.08361*, 2020.
- [6] N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean, “Outrageously large neural networks: The sparsely-gated mixture-of-experts layer,” *arXiv preprint arXiv:1701.06538*, 2017.
- [7] W. Fedus, B. Zoph, and N. Shazeer, “Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity,” *Journal of Machine Learning Research*, vol. 23, no. 120, pp. 1–39, 2022.
- [8] J. Achiam, S. Adler, S. Agarwal, L. Ahmad, I. Akkaya, F. L. Aleman, D. Almeida, J. Altschmidt, S. Altman, S. Anadkat *et al.*, “Gpt-4 technical report,” *arXiv preprint arXiv:2303.08774*, 2023.
- [9] D. Dai, C. Deng, C. Zhao, R. Xu, H. Gao, D. Chen, J. Li, W. Zeng, X. Yu, Y. Wu *et al.*, “Deepseekmoe: Towards ultimate expert specialization in mixture-of-experts language models,” *arXiv preprint arXiv:2401.06066*, 2024.
- [10] Q. Team, “Qwen1.5-moe: Matching 7b model performance with 1/3 activated parameters,” February 2024. [Online]. Available: <https://qwenlm.github.io/blog/qwen-moe/>
- [11] A. Yang, A. Li, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Gao, C. Huang, C. Lv *et al.*, “Qwen3 technical report,” *arXiv preprint arXiv:2505.09388*, 2025.
- [12] OpenAI, “Introducing gpt-oss,” 2025. [Online]. Available: <https://openai.com/index/introducing-gpt-oss/>
- [13] M. Abdin, J. Aneja, H. Awadalla, A. Awadallah, A. A. Awan, N. Bach, A. Bahree, A. Bakhtiari, J. Bao, H. Behl, and *et al.*, “Phi-3 technical report: A highly capable language model locally on your phone,” 2024. [Online]. Available: <https://arxiv.org/abs/2404.14219>
- [14] Mistral, “Mistral of experts,” 2023. [Online]. Available: <https://mistral.ai/news/mistral-of-experts>
- [15] T. H. Team, “Hunyuan a13b technical report,” 2025. [Online]. Available: https://github.com/Tencent-Hunyuan/Hunyuan-A13B/blob/main/report/Hunyuan_A13B_Technical_Report.pdf
- [16] Y. Tang, X. Li, F. Liu, W. Guo, H. Zhou, Y. Wang, K. Han, X. Yu, J. Li, H. Zang *et al.*, “Pangu pro moe: Mixture of grouped experts for efficient sparsity,” *arXiv preprint arXiv:2505.21411*, 2025.
- [17] L. Weidinger, J. Uesato, M. Rauh, C. Griffin, P.-S. Huang, J. Mellor, A. Glaese, M. Cheng, B. Balle, A. Kasirzadeh *et al.*, “Taxonomy of risks posed by language models,” in *Proceedings of the 2022 ACM conference on fairness, accountability, and transparency*, 2022, pp. 214–229.
- [18] A. Cloud, M. Le, J. Chua, J. Betley, A. Szyber-Betley, J. Hilton, S. Marks, and O. Evans, “Subliminal learning: Language models transmit behavioral traits via hidden signals in data,” *arXiv preprint arXiv:2507.14805*, 2025.
- [19] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray *et al.*, “Training language models to follow instructions with human feedback,” *Advances in neural information processing systems*, vol. 35, pp. 27 730–27 744, 2022.
- [20] R. Rafailov, A. Sharma, E. Mitchell, C. D. Manning, S. Ermon, and C. Finn, “Direct preference optimization: Your language model is secretly a reward model,” *Advances in neural information processing systems*, vol. 36, pp. 53 728–53 741, 2023.
- [21] E. Perez, S. Huang, F. Song, T. Cai, R. Ring, J. Aslanides, A. Glaese, N. McAleese, and G. Irving, “Red teaming language models with language models,” *arXiv preprint arXiv:2202.03286*, 2022.
- [22] A. Wei, N. Haghtalab, and J. Steinhardt, “Jailbroken: How does llm safety training fail?” *Advances in Neural Information Processing Systems*, vol. 36, pp. 80 079–80 110, 2023.
- [23] X. Shen, Z. Chen, M. Backes, Y. Shen, and Y. Zhang, ““do anything now”: Characterizing and evaluating in-the-wild jailbreak prompts on large language models,” in *Proceedings of the 2024 on ACM SIGSAC Conference on Computer and Communications Security*, 2024, pp. 1671–1685.
- [24] S. Li, L. Yao, L. Zhang, and Y. Li, “Safety layers in aligned large language models: The key to llm security,” *arXiv preprint arXiv:2408.17003*, 2024.

- [25] J. Chen, X. Wang, Z. Yao, Y. Bai, L. Hou, and J. Li, “Finding safety neurons in large language models,” *arXiv preprint arXiv:2406.14144*, 2024.
- [26] A. Zou, Z. Wang, N. Carlini, M. Nasr, J. Z. Kolter, and M. Fredrikson, “Universal and transferable adversarial attacks on aligned language models,” *arXiv preprint arXiv:2307.15043*, 2023.
- [27] Q. Wang, Q. Pang, X. Lin, S. Wang, and D. Wu, “Badmoe: Backdooring mixture-of-experts llms via optimizing routing triggers and infecting dormant experts,” *arXiv preprint arXiv:2504.18598*, 2025.
- [28] Z. Lai, M. Liao, D. Xu, Z. Zhao, Z. Yuan, C. Fan, J. Li, and B. Wu, “Safex: Analyzing vulnerabilities of moe-based llms via stable safety-critical expert identification,” *arXiv preprint arXiv:2506.17368*, 2025.
- [29] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin, “Attention is all you need,” *Advances in neural information processing systems*, vol. 30, 2017.
- [30] W. Cai, J. Jiang, F. Wang, J. Tang, S. Kim, and J. Huang, “A survey on mixture of experts in large language models,” *IEEE Transactions on Knowledge and Data Engineering*, 2025.
- [31] T. Krauß, H. Dashtbani, and A. Dmitrienko, “Twinbreak: Jailbreaking llm security alignments based on twin prompts,” *arXiv preprint arXiv:2506.07596*, 2025.
- [32] LLM-LAT, “harmful-dataset,” 2024. [Online]. Available: <https://huggingface.co/datasets/LLM-LAT/harmful-dataset>
- [33] R. Bhardwaj, D. D. Anh, and S. Poria, “Language models are homer simpson! safety re-alignment of fine-tuned language models through task arithmetic,” 2024.
- [34] R. Bhardwaj and S. Poria, “Red-teaming large language models using chain of utterances for safety-alignment,” 2023.
- [35] W. Luo, S. Ma, X. Liu, X. Guo, and C. Xiao, “Jailbreakv-28k: A benchmark for assessing the robustness of multimodal large language models against jailbreak attacks,” 2024.
- [36] T. Kwiatkowski, J. Palomaki, O. Redfield, M. Collins, A. Parikh, C. Alberti, D. Epstein, I. Polosukhin, J. Devlin, K. Lee *et al.*, “Natural questions: a benchmark for question answering research,” *Transactions of the Association for Computational Linguistics*, vol. 7, pp. 453–466, 2019.
- [37] M. Geva, A. Caciularu, K. R. Wang, and Y. Goldberg, “Transformer feed-forward layers build predictions by promoting concepts in the vocabulary space,” *arXiv preprint arXiv:2203.14680*, 2022.
- [38] H. J. Davies, “Decoding specialised feature neurons in llms with the final projection layer,” *arXiv preprint arXiv:2501.02688*, 2025.
- [39] H. Inan, K. Upasani, J. Chi, R. Rungta, K. Iyer, Y. Mao, M. Tontchev, Q. Hu, B. Fuller, D. Testuggine *et al.*, “Llama guard: Llm-based input-output safeguard for human-ai conversations,” *arXiv preprint arXiv:2312.06674*, 2023.
- [40] H. Zhao, C. Yuan, F. Huang, X. Hu, Y. Zhang, A. Yang, B. Yu, D. Liu, J. Zhou, J. Lin *et al.*, “Qwen3guard technical report,” *arXiv preprint arXiv:2510.14276*, 2025.
- [41] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman, “Glue: A multi-task benchmark and analysis platform for natural language understanding,” *arXiv preprint arXiv:1804.07461*, 2018.
- [42] K. Sakaguchi, R. L. Bras, C. Bhagavatula, and Y. Choi, “Winogrande: An adversarial winograd schema challenge at scale,” *Communications of the ACM*, vol. 64, no. 9, pp. 99–106, 2021.
- [43] T. Mihaylov, P. Clark, T. Khot, and A. Sabharwal, “Can a suit of armor conduct electricity? a new dataset for open book question answering,” *arXiv preprint arXiv:1809.02789*, 2018.
- [44] P. Clark, I. Cowhey, O. Etzioni, T. Khot, A. Sabharwal, C. Schoenick, and O. Tafjord, “Think you have solved question answering? try arc, the ai2 reasoning challenge,” *arXiv preprint arXiv:1803.05457*, 2018.
- [45] A. Souly, Q. Lu, D. Bowen, T. Trinh, E. Hsieh, S. Pandey, P. Abbeel, J. Svegliato, S. Emmons, O. Watkins, and S. Toyer, “A strongreject for empty jailbreaks,” 2024.
- [46] X. Qi, Y. Zeng, T. Xie, P.-Y. Chen, R. Jia, P. Mittal, and P. Henderson, “Fine-tuning aligned language models compromises safety, even when users do not intend to!” *arXiv preprint arXiv:2310.03693*, 2023.
- [47] Z. Wu, X. Chen, Z. Pan, X. Liu, W. Liu, D. Dai, H. Gao, Y. Ma, C. Wu, B. Wang *et al.*, “Deepseek-vl2: Mixture-of-experts vision-language models for advanced multimodal understanding,” *arXiv preprint arXiv:2412.10302*, 2024.
- [48] K. Team, A. Du, B. Yin, B. Xing, B. Qu, B. Wang, C. Chen, C. Zhang, C. Du, C. Wei *et al.*, “Kimi-vl technical report,” *arXiv preprint arXiv:2504.07491*, 2025.

- [49] Y. Gong, D. Ran, J. Liu, C. Wang, T. Cong, A. Wang, S. Duan, and X. Wang, “Figstep: Jailbreaking large vision-language models via typographic visual prompts,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 39, no. 22, 2025, pp. 23 951–23 959.
- [50] F. Weng, Y. Xu, C. Fu, and W. Wang, “Mmj-bench: A comprehensive study on jailbreak attacks and defenses for multimodal large language models,” *arXiv preprint arXiv:2408.08464*, 2024.
- [51] E. Nelson, N. Neel, O. Catherine, H. Tom, J. Nicholas, M. Ben, A. Amanda, B. Yuntao, C. Anna, C. Tom *et al.*, “A mathematical framework for transformer circuits,” *Transformer Circuits Thread*, 2021.
- [52] L. Wu, S. Behrouzi, M. Rostami, M. Thang, S. Picek, and A.-R. Sadeghi, “Neurostrike: Neuron-level attacks on aligned llms,” in *NDSS*, 2026.
- [53] F. Perez and I. Ribeiro, “Ignore previous prompt: Attack techniques for language models,” *arXiv preprint arXiv:2211.09527*, 2022.
- [54] S. Jiang, X. Chen, and R. Tang, “Prompt packer: Deceiving llms through compositional instruction with hidden attacks,” *arXiv preprint arXiv:2310.10077*, 2023.
- [55] X. Li, Z. Zhou, J. Zhu, J. Yao, T. Liu, and B. Han, “Deepinception: Hypnotize large language model to be jailbreaker,” *arXiv preprint arXiv:2311.03191*, 2023.
- [56] J. Yu, X. Lin, Z. Yu, and X. Xing, “Gptfuzzer: Red teaming large language models with auto-generated jailbreak prompts,” *arXiv preprint arXiv:2309.10253*, 2023.
- [57] J. Ebrahimi, A. Rao, D. Lowd, and D. Dou, “Hotflip: White-box adversarial examples for text classification,” *arXiv preprint arXiv:1712.06751*, 2017.
- [58] Y. Wen, N. Jain, J. Kirchenbauer, M. Goldblum, J. Geiping, and T. Goldstein, “Hard prompts made easy: Gradient-based discrete optimization for prompt tuning and discovery,” *Advances in Neural Information Processing Systems*, vol. 36, pp. 51 008–51 025, 2023.
- [59] Y. Zhou, A. I. Muresanu, Z. Han, K. Paster, S. Pitis, H. Chan, and J. Ba, “Large language models are human-level prompt engineers,” in *The Eleventh International Conference on Learning Representations*, 2022.
- [60] P. Chao, A. Robey, E. Dobriban, H. Hassani, G. J. Pappas, and E. Wong, “Jailbreaking black box large language models in twenty queries,” *arXiv preprint arXiv:2310.08419*, 2023.
- [61] Z. Chang, M. Li, Y. Liu, J. Wang, Q. Wang, and Y. Liu, “Play guessing game with llm: Indirect jailbreak attack with implicit clues,” *arXiv preprint arXiv:2402.09091*, 2024.
- [62] A. Kádár, G. Chrupała, and A. Alishahi, “Representation of linguistic form and function in recurrent neural networks,” *Computational Linguistics*, vol. 43, no. 4, pp. 761–780, 2017.
- [63] Y. Lakretz, G. Kruszewski, T. Desbordes, D. Hupkes, S. Dehaene, and M. Baroni, “The emergence of number and syntax units in lstm language models,” *arXiv preprint arXiv:1903.07435*, 2019.
- [64] Anthropic, “Scaling monosemanticity: Extracting interpretable features from claude 3 sonnet,” 2024. [Online]. Available: <https://transformer-circuits.pub/2024/scaling-monosemanticity/>
- [65] B. Wei, K. Huang, Y. Huang, T. Xie, X. Qi, M. Xia, P. Mittal, M. Wang, and P. Henderson, “Assessing the brittleness of safety alignment via pruning and low-rank modifications,” in *Proceedings of the 41st International Conference on Machine Learning*, 2024, pp. 52 588–52 610.
- [66] J. Hayes, I. Shumailov, and I. Yona, “Buffer overflow in mixture of experts,” *arXiv preprint arXiv:2402.05526*, 2024.
- [67] I. Yona, I. Shumailov, J. Hayes, and N. Carlini, “Stealing user prompts from mixture of experts,” *arXiv preprint arXiv:2410.22884*, 2024.

Appendix

A Sparse Expert Ablation

As stated in Section 2.2, the router sparsely activates the top-K experts in each MoE layer. Since this selection strongly influences the model’s response to different prompts, we treat K as a key hyperparameter. We conduct sparse expert ablation by progressively removing the top-0.5K, top-K, and top-2K experts per layer, ranked by their malicious utility scores. As shown in Table 11, this targeted ablation consistently increases the Attack Success Rate (ASR) across all evaluated MoE LLMs, demonstrating the safety relevance of the most active experts during harmful input processing. Overall, the average ASR increases from 19.3% (top-0.5K) to 45.9% (top-2K), confirming that experts with high malicious utility scores play a central role in mediating refusal behavior.

Target Model	Top-0.5K	Top-K	Top-2K
GPT-OSS-20B	15.0%	32.9%	73.8%
Qwen3-30B-A3B-Instruct-2507	4.8%	26.5%	39.6%
Phi-3.5-MoE-Instruct	0.0%	47.3%	58.8%
Mixtral-8x7B-Instruct-v0.1	56.9%	21.7%	33.5%
Qwen1.5-MoE-A2.7B-Chat	10.9%	10.5%	15.7%
Deepseek-MoE-16b-Chat	47.9%	70.0%	57.2%
Hunyuan-A13B-Instruct	11.8%	14.1%	71.9%
OpenPangu-Pro-MoE-72B	6.7%	8.6%	16.6%
<i>Average</i>	<i>19.3%</i>	<i>29.0%</i>	<i>45.9%</i>

Table 11: ASR with Sparse Expert Ablation.

Notably, the ASR gains are not uniform across models. For example, GPT-OSS-20B shows a dramatic rise from 15.0% (top-0.5K) to 73.8% (top-2K), highlighting a strong concentration of safety behavior in a relatively small subset of experts. Similarly, Hunyuan-A13B-Instruct and Phi-3.5-MoE-Instruct also show steep increases at the top-2K level, reaching 71.9% and 58.8%, respectively. These results indicate that safety alignment in MoE models is not widely distributed, but rather localized to specific, highly utilized experts. Interestingly, Mixtral-8x7B-Instruct-v0.1 exhibits an unusually high ASR even at the top-0.5K level (56.9%), but a more modest increase across the broader ablation range, suggesting that a small number of experts may disproportionately influence safety outcomes in that model. In contrast, Qwen1.5-MoE-A2.7B-Chat and OpenPangu-Pro-MoE-72B show more gradual increases, which may reflect a more diffused or redundant safety structure.

B Safety Neuron Suppression Strength

To understand how strongly safety neurons contribute to refusal behavior, we vary the suppression strength applied to identified safety neurons and measure the resulting ASR.

Specifically, we evaluate partial activation suppression at 35% and 65%, and compare them against full suppression (100%).

Target Model	35%	65%	100%
GPT-OSS-20B	9.9%	68.6%	80.2%
Qwen3-30B-A3B-Instruct-2507	2.2%	22.6%	56.9%
Phi-3.5-MoE-Instruct	1.5%	11.5%	56.5%
Mixtral-8x7B-Instruct-v0.1	36.7%	54.6%	64.2%
Qwen1.5-MoE-A2.7B-Chat	25.6%	42.5%	57.8%
Deepseek-MoE-16b-Chat	25.2%	29.4%	53.6%
Hunyuan-A13B-Instruct	42.2%	66.1%	76.9%
OpenPangu-Pro-MoE-72B	35.8%	64.9%	73.1%
<i>Average</i>	<i>22.4%</i>	<i>45.0%</i>	<i>64.9%</i>

Table 12: ASR with different numbers of safety experts.

As shown in Table 12, ASR increases monotonically with stronger suppression. On average, partial suppression at 35% and 65% yields ASRs of 22.4% and 45.0%, respectively, both substantially higher than the baseline ASR of 7.4%, yet still noticeably lower than the 64.9% achieved with full suppression. This trend is consistent across all evaluated MoE models, despite their architectural and routing differences.

These results indicate that safety behavior emerges from the collective contribution of multiple safety neurons. Partial suppression weakens refusal behavior proportionally, while full suppression is required to reliably disable it. The smooth degradation in safety performance further suggests a strong causal relationship between the activation magnitude of safety neurons and the model’s refusal behavior, rather than an all-or-nothing effect.

C The Importance of Safety Neurons

To verify that the observed effects are specific to safety neurons rather than a byproduct of indiscriminate neuron disruption, we compare targeted safety neuron suppression against random neuron pruning. We conduct two random baselines: R1, which prunes an equal number of neurons randomly selected from all experts, and R2, which randomly prunes neurons within the safety expert subsets. Table 13 shows that

Target Model	R1	R2	GateBreaker
GPT-OSS-20B	1.3%	1.6%	80.2%
Qwen3-30B-A3B-Instruct-2507	0.6%	0.6%	56.9%
Phi-3.5-MoE-Instruct	0.3%	0.6%	56.5%
Mixtral-8x7B-Instruct-v0.1	12.8%	16.3%	64.2%
Qwen1.5-MoE-A2.7B-Chat	5.1%	7.3%	57.8%
Deepseek-MoE-16b-Chat	23.3%	18.5%	53.6%
Hunyuan-A13B-Instruct	12.5%	28.4%	76.9%
OpenPangu-Pro-MoE-72B	6.6%	21.1%	73.1%
<i>Average</i>	<i>7.8%</i>	<i>11.8%</i>	<i>64.9%</i>

Table 13: ASR with random/safety neuron pruning.

random neuron pruning results in consistently low ASR, with averages of 7.8% and 11.8%, closely matching the baseline behavior. In contrast, GateBreaker achieves a significantly higher average ASR of 64.9%, with improvements exceeding $5\times$ to $10\times$ across most models. Indeed, the ineffectiveness of random pruning confirms that the observed jailbreak success is not due to reduced model capacity or degraded fluency, but rather to the targeted disruption of neurons that encode safety-relevant representations.